

Optimization of the C program summary

The general approach to optimize the code was done in a few steps.

Step 1:

- convert derivations of constants from other constants, e.g. $a = 3.3$ $b = \log(a)$, to a constant value. This is only noticeable when optimization is turned off.
- convert $\text{pow}(x, \text{integer})$ to a series of multiplies. This replaces a several dozen operations to only a few multiplies.

Step 2:

- convert $\text{pow}(x, \text{float } y)$ to $\exp(y * \log(x))$. The exponential function is faster than the power function. In this code there are multiple calls to the power function of the same significand $\text{pow}(\text{fpt}, \text{blah})$. If we precompute the value of $\log(\text{fpt})$, there are big savings.
- convert the loops over floating point values in main functions to ones over integers.
- Multithread the code with openmp so several points can be done at once. We can expect a speedup proportional to the number of cores.

I ran a series of timed tests (using bash time) on each version of the code with a linux desktop running Ubuntu 13 with the gcc-13 compiler. The CPU was an Intel(R) Core(TM) i7-3820 CPU @ 3.60GHz with 4 cores and hyperthreading purchased in 2011. All times were collected using the bash time utility and the speedups are relative to the original code.

time(secs)			
version	no optimization	-O	-O2 -ffast-math
original	436.83	234.85	0.934
change 1	191.64	171.15	0.938
change 2	121.27	109.5	0.002
change 2 + mp	27.92	22.15	0.001

speed up			
original	1	1	1
change 1	2.279430182	1.372188139	0.9957356077
change 2	3.602127484	2.144748858	467
change 2 + mp	15.64577364	10.6027088	934

Profiles obtained with gprof utility are available in the github repo. Note that the openmp speedup is roughly 4 times that of the unthreaded code which is to be expected for a CPU with 4 cores. Overall, the optimizations were fairly effective in reducing the runtime.